

```
In [1]: import os, random, copy
from pathlib import Path

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from PIL import Image
from tqdm.notebook import tqdm

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader
from torchvision import datasets, transforms, models
from torchvision.datasets import ImageFolder
from sklearn.metrics import classification_report, confusion_matrix

import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: # — Reproducibility —
def set_seed(seed=42):
    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    torch.cuda.manual_seed_all(seed)
    torch.backends.cudnn.deterministic = True
    torch.backends.cudnn.benchmark = False

set_seed(42)
DEVICE = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f'Using device: {DEVICE}')

# — Paths (matches your existing EDA setup) —
BASE_DIR = Path('..').resolve()
DATA_DIR = BASE_DIR / 'data' / 'raw'

splits = {
    'train': DATA_DIR / 'train',
    'validate': DATA_DIR / 'validate',
    'test': DATA_DIR / 'test',
}
```

Using device: cpu

1. EDA

```
In [3]: def win_safe_loader(path: str):
    if os.name == 'nt' and not path.startswith('\\?\\'):
        path = '\\?\\' + path
    with open(path, 'rb') as f:
        img = Image.open(f)
        return img.convert('RGB')

class SafeImageFolder(ImageFolder):
    def __init__(self, root, transform=None, **kwargs):
```

```
super().__init__(root=root, transform=transform,
                 loader=win_safe_loader, **kwargs)
```

```
In [4]: # 1. PATH SETUP (matches your folder layout)
# Your notebook is in: project/notebooks/ So go UP ONE LEVEL to reach the project
# In computer-speak, .. means "go up one level." It refers to the parent folder
#.resolve() forces Python to Look at the actual address on your hard drive. It takes us
# into an Absolute Path (like C:\Users\glori\Documents\...).
BASE_DIR = Path("..").resolve()
DATA_DIR = BASE_DIR / "data" / "raw"

# A dictionary mapping each split name to its folder path. Using a dict instead
splits = {
    "train": DATA_DIR / "train",
    "validate": DATA_DIR / "validate",
    "test": DATA_DIR / "test"
}

# classes drives everything – the counting, totalling, and melting all reference
classes = ["edible", "poisonous"]
# a list of dictionaries. Each dictionary looked like {"split": "train", "edible": ...
valid_extensions = {".jpg", ".jpeg", ".png"}
rows = []

for split, dir_path in splits.items():
    # Starts a new dictionary for this split. After the inner loop it will look like:
    # python{"split": "train", "edible": 240, "poisonous": 180, "total": 420}
    row = {"split": split}

    # .glob("*") returns a generator of Path objects, one for each file. Example
    # sum(1 for ...) → counts the matches by adding 1 for each valid file (more)
    for cls in classes:
        row[cls] = sum(1 for f in (dir_path / cls).glob("*") if f.suffix.lower() == cls.lower())
        # So after the inner loop, row looks like:
        # python{"split": "train", "edible": 240, "poisonous": 180}

    # Adds up all class counts. Using classes here means if you add a third class
    row["total"] = sum(row[c] for c in classes)
    # Adds the completed row dict to the list
    rows.append(row)

# a constructor from the Pandas Library. It takes that list rows and automatically creates a DataFrame
df = pd.DataFrame(rows)
# When you type the name of a variable as the very last line in a code cell (esp
# Without this line, the DataFrame would be created silently in the background.
df
```

```
Out[4]:
```

	split	edible	poisonous	total
0	train	1200	1056	2256
1	validate	150	132	282
2	test	150	132	282

```
In [5]: #BEFORE melt (wide)
#AFTER melt (long):
#split      mushroom type   number of samples
#train      edible          1200
```

```


#train      poisonous      1056
# standard pandas melt parameters/parameter names of the .melt() function. always
# your DataFrame does NOT need to have columns with these names. they are just there
df_melted = df.melt(
    id_vars='split', # Keep `split` column as-is. columns that should stay as they are
    value_vars=classes, # columns that should be unpivoted into rows, turned into columns
    var_name='mushroom type', # The new column that holds the old column names
    value_name='number of samples' # The new column that holds the actual counts
)

plt.figure(figsize=(10, 6))
ax = sns.barplot(data=df_melted, x='split', y='number of samples', hue='mushroom type')

for container in ax.containers:
    ax.bar_label(container)

plt.legend(title="mushroom type")
plt.title('mushroom class distribution per split (grouped bar chart)')
plt.xlabel('dataset split')
plt.ylabel('number of samples')
plt.show()

```



In [8]:

```

# set up the plot
fig, axes = plt.subplots(2, 3, figsize=(15, 10))
fig.suptitle('sample images: each class per split', fontsize=16, y=1.02)

# flatten axes for easier indexing
axes_flat = axes.flatten()
# counter for subplot
plot_idx = 0

# for each split and class, display a random sample image
for split, dir_path in splits.items():
    for cls in classes:
        # get all image files in the class folder
        class_path = dir_path / cls
        images = [f for f in class_path.glob("*") if f.suffix.lower() in valid_extensions]

```

```

# select a random image
if images:
    random_image = random.choice(images)
    # Load and display the image. use the safe function here
    img = win_safe_loader(str(random_image))
    axes_flat[plot_idx].imshow(img)
    axes_flat[plot_idx].set_title(f'{split} - {cls}', fontsize=12)
    axes_flat[plot_idx].axis('off')
else:
    # Handle case with no images
    axes_flat[plot_idx].text(0.5, 0.5, f'No images\nin {split}/{cls}',
                           ha='center', va='center')
    axes_flat[plot_idx].set_title(f'{split} - {cls}', fontsize=12)
    axes_flat[plot_idx].axis('off')

plot_idx += 1

plt.tight_layout()
plt.show()

```

sample images: each class per split



2. Shared Training Engine

One reusable function used for **both** the Baseline CNN and Transfer Learning.

What this combines from all versions:

- `tqdm` progress bars with live loss/acc (Doc 5)
- `copy.deepcopy` best-model checkpointing — restores best weights at end (Doc 5)
- `finetune_epoch` auto-switch inside the loop (Doc 5)
- Differential learning rates for backbone vs head during fine-tuning (my improvement)
- 85% target line + fine-tune vertical line on accuracy plot (Docs 4 & 5)

```
In [9]: def train_model(model, train_loader, val_loader, criterion, optimizer,
                    scheduler=None, num_epochs=15, finetune_epoch=None,
                    finetune_layers=('layer4', 'fc'), finetune_lr=1e-4):
    """
    Universal training loop.

    Parameters
    -----
    finetune_epoch : int or None -- epoch index at which to unfreeze layers
    finetune_layers : tuple -- layer name substrings to unfreeze
    finetune_lr : float -- LR for backbone; head gets 10x this value
    """
    history = {'train_loss': [], 'val_loss': [],
               'train_acc': [], 'val_acc': []}
    best_val_acc = 0.0
    best_weights = copy.deepcopy(model.state_dict())
    best_preds, best_labels = [], []

    for epoch in range(num_epochs):

        # — Phase switch: unfreeze Layers & apply differential LRs ————
        if finetune_epoch and epoch == finetune_epoch:
            print(f'\n{"*55}')
            print(f' Epoch {epoch+1}: SWITCHING TO FINE-TUNE – unfreezing {fine}
print(f'{"\n*55\n"}')
            for name, param in model.named_parameters():
                if any(layer in name for layer in finetune_layers):
                    param.requires_grad = True
        # Differential LRs: backbone gets small LR to protect pretrained features
        # head gets 10x larger LR because it is newly initialised
        backbone_params = [p for n, p in model.named_parameters()
                            if p.requires_grad and 'fc' not in n]
        head_params = [p for n, p in model.named_parameters()
                       if p.requires_grad and 'fc' in n]
        optimizer = optim.Adam([
            {'params': backbone_params, 'lr': finetune_lr},
            {'params': head_params, 'lr': finetune_lr * 10},
        ])

        # — Training phase ————
        model.train()
        running_loss, correct, total = 0.0, 0, 0

        train_bar = tqdm(train_loader, desc=f'Epoch {epoch+1}/{num_epochs} [{len(train_loader)}]')
        for imgs, labels in train_bar:
            imgs, labels = imgs.to(DEVICE), labels.to(DEVICE)
            optimizer.zero_grad()
            outputs = model(imgs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            running_loss += loss.item() * imgs.size(0)
            _, preds = outputs.max(1)
            correct += preds.eq(labels).sum().item()
            total += imgs.size(0)
            train_bar.set_postfix(loss=f'{loss.item():.4f}', acc=f'{100*correct/total:.1f}%')
```

```

epoch_train_loss = running_loss / total
epoch_train_acc = 100 * correct / total

# — Validation phase ——————
model.eval()
val_loss, correct, total = 0.0, 0, 0
epoch_preds, epoch_labels = [], []

with torch.no_grad():
    val_bar = tqdm(val_loader, desc=f'Epoch {epoch+1:02d}/{num_epochs} [')
    for imgs, labels in val_bar:
        imgs, labels = imgs.to(DEVICE), labels.to(DEVICE)
        outputs = model(imgs)
        loss = criterion(outputs, labels)

        val_loss += loss.item() * imgs.size(0)
        _, preds = outputs.max(1)
        correct += preds.eq(labels).sum().item()
        total += imgs.size(0)
        epoch_preds.extend(preds.cpu().numpy())
        epoch_labels.extend(labels.cpu().numpy())
        val_bar.set_postfix(loss=f'{loss.item():.4f}', acc=f'{100*correct/total:.1f}%')

epoch_val_loss = val_loss / total
epoch_val_acc = 100 * correct / total

if scheduler:
    scheduler.step()

history['train_loss'].append(epoch_train_loss)
history['val_loss'].append(epoch_val_loss)
history['train_acc'].append(epoch_train_acc)
history['val_acc'].append(epoch_val_acc)

print(f'Epoch {epoch+1:02d}/{num_epochs} '
      f'Train Loss: {epoch_train_loss:.4f} Acc: {epoch_train_acc:.2f}%
      f'Val Loss: {epoch_val_loss:.4f} Acc: {epoch_val_acc:.2f}%)')

# — Best-model checkpoint ——————
if epoch_val_acc > best_val_acc:
    best_val_acc = epoch_val_acc
    best_weights = copy.deepcopy(model.state_dict())
    best_preds, best_labels = epoch_preds[:], epoch_labels[:]
    print(f' ★ New best model saved (Val Acc: {best_val_acc:.2f}%)')

model.load_state_dict(best_weights) # restore best weights
return model, history, best_preds, best_labels

def plot_curves(history, title='Training Curves', finetune_epoch=None):
    """Loss & accuracy curves with 85% target line and optional fine-tune marker
    epochs = range(1, len(history['train_loss']) + 1)
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 5))

    ax1.plot(epochs, history['train_loss'], 'o-', label='Train Loss')
    ax1.plot(epochs, history['val_loss'], 'o-', label='Val Loss')
    ax1.set(title='Loss', xlabel='Epoch', ylabel='Loss')
    ax1.legend(); ax1.grid(True)

```

```

ax2.plot(epochs, history['train_acc'], 'o-', label='Train Acc')
ax2.plot(epochs, history['val_acc'], 'o-', label='Val Acc')
ax2.axhline(y=85, color='red', linestyle='--', label='85% target')      # from
if finetune_epoch is not None:
    ax2.axvline(x=finetune_epoch + 1, color='orange',
                 linestyle='--', label='Fine-tune start')                # from
ax2.set(title='Accuracy (%)', xlabel='Epoch', ylabel='Accuracy (%)')
ax2.legend(); ax2.grid(True)

fig.suptitle(title, fontsize=13)
plt.tight_layout()
plt.show()

```

3. Baseline CNN

Target: ~60–70% val accuracy — just proving the problem is learnable.

```

In [ ]: # — Transforms (128x128, light augmentation) —————
baseline_train_tf = transforms.Compose([
    transforms.Resize((128, 128)),
    transforms.RandomHorizontalFlip(p=0.5),
    transforms.RandomRotation(10),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225]),
])
baseline_val_tf = transforms.Compose([
    transforms.Resize((128, 128)),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225]),
])

# uses train and validate splits during training
train_ds_base = SafeImageFolder(root=str(splits['train']), transform=baseline
val_ds_base = SafeImageFolder(root=str(splits['validate']), transform=baseline
print(f'Classes: {train_ds_base.classes}')
print(f'Train: {len(train_ds_base)} | Val: {len(val_ds_base)}')

train_dl_base = DataLoader(train_ds_base, batch_size=32, shuffle=True, num_workers=4)
val_dl_base = DataLoader(val_ds_base, batch_size=32, shuffle=False, num_workers=4)

```

Classes: ['edible', 'poisonous']
Train: 2,256 | Val: 282

```

In [11]: class BaselineCNN(nn.Module):
    """
    3 conv blocks: Conv -> BatchNorm -> ReLU -> MaxPool
    128x128 -> 64x64 -> 32x32 -> 16x16 feature maps, then FC head.
    """

    def __init__(self, num_classes=2):
        super().__init__()
        self.features = nn.Sequential(
            nn.Conv2d(3, 32, 3, padding=1), nn.BatchNorm2d(32), nn.ReLU(), nn.MaxPool2d(2),
            nn.Conv2d(32, 64, 3, padding=1), nn.BatchNorm2d(64), nn.ReLU(), nn.MaxPool2d(2),
            nn.Conv2d(64, 128, 3, padding=1), nn.BatchNorm2d(128), nn.ReLU(), nn.MaxPool2d(2)
        )
        self.classifier = nn.Sequential(
            nn.Flatten(),
            nn.Dropout(0.3), nn.Linear(128*16*16, 256), nn.ReLU(),
            nn.Linear(256, num_classes)
        )

```

```

        nn.Dropout(0.3), nn.Linear(256, num_classes),
    )

    def forward(self, x):
        return self.classifier(self.features(x))

baseline_model = BaselineCNN().to(DEVICE)

criterion_b = nn.CrossEntropyLoss()
optimizer_b = optim.Adam(baseline_model.parameters(), lr=1e-3)
scheduler_b = optim.lr_scheduler.StepLR(optimizer_b, step_size=5, gamma=0.5)

baseline_model, history_base, _, _ = train_model(
    baseline_model, train_dl_base, val_dl_base,
    criterion_b, optimizer_b, scheduler_b,
    num_epochs=15
)

plot_curves(history_base, title='Baseline CNN - Training Curves')
print(f'Baseline CNN - Best Val Accuracy: {max(history_base["val_acc"]):.2f}%')

```


Epoch 01/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 01/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 01/15 Train Loss: 1.9924 Acc: 55.27% | Val Loss: 0.6309 Acc: 67.02%
★ New best model saved (Val Acc: 67.02%)
Epoch 02/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 02/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 02/15 Train Loss: 0.6485 Acc: 60.73% | Val Loss: 0.6179 Acc: 66.31%
Epoch 03/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 03/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 03/15 Train Loss: 0.6441 Acc: 60.55% | Val Loss: 0.6258 Acc: 68.79%
★ New best model saved (Val Acc: 68.79%)
Epoch 04/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 04/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 04/15 Train Loss: 0.6463 Acc: 59.35% | Val Loss: 0.6203 Acc: 67.73%
Epoch 05/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 05/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 05/15 Train Loss: 0.6424 Acc: 61.66% | Val Loss: 0.5949 Acc: 69.15%
★ New best model saved (Val Acc: 69.15%)
Epoch 06/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 06/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 06/15 Train Loss: 0.6252 Acc: 63.56% | Val Loss: 0.6076 Acc: 65.96%
Epoch 07/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 07/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 07/15 Train Loss: 0.6180 Acc: 63.21% | Val Loss: 0.6008 Acc: 68.79%
Epoch 08/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 08/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 08/15 Train Loss: 0.6079 Acc: 65.47% | Val Loss: 0.5957 Acc: 67.73%
Epoch 09/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 09/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 09/15 Train Loss: 0.6091 Acc: 65.56% | Val Loss: 0.5880 Acc: 66.67%
Epoch 10/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 10/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 10/15 Train Loss: 0.5936 Acc: 66.13% | Val Loss: 0.6041 Acc: 65.25%
Epoch 11/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 11/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 11/15 Train Loss: 0.5811 Acc: 67.82% | Val Loss: 0.5738 Acc: 68.79%
Epoch 12/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 12/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]

```

Epoch 12/15 Train Loss: 0.5718 Acc: 67.69% | Val Loss: 0.5840 Acc: 69.50%
★ New best model saved (Val Acc: 69.50%)
Epoch 13/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 13/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 13/15 Train Loss: 0.5661 Acc: 67.60% | Val Loss: 0.5582 Acc: 69.50%
Epoch 14/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 14/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 14/15 Train Loss: 0.5593 Acc: 69.68% | Val Loss: 0.5484 Acc: 72.70%
★ New best model saved (Val Acc: 72.70%)
Epoch 15/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 15/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 15/15 Train Loss: 0.5488 Acc: 70.08% | Val Loss: 0.5877 Acc: 68.09%

```

Baseline CNN — Training Curves



Baseline CNN — Best Val Accuracy: 72.70%

4. Transfer Learning — ResNet18

Target: >85% val accuracy using ImageNet pretrained weights + fine-tuning.

```

In [12]: # — Transforms (224x224, heavy augmentation) —————
IMAGENET_MEAN = [0.485, 0.456, 0.406]
IMAGENET_STD = [0.229, 0.224, 0.225]

tl_train_tf = transforms.Compose([
    transforms.Resize((256, 256)),
    transforms.RandomResizedCrop(224, scale=(0.7, 1.0)),
    transforms.RandomHorizontalFlip(p=0.5),
    transforms.RandomRotation(15),
    transforms.ColorJitter(brightness=0.3, contrast=0.3, saturation=0.2, hue=0.1),
    transforms.ToTensor(),
    transforms.Normalize(IMAGENET_MEAN, IMAGENET_STD),
])
tl_val_tf = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(IMAGENET_MEAN, IMAGENET_STD),
])

train_ds_tl = SafeImageFolder(root=str(splits['train']), transform=tl_train_tf)
val_ds_tl = SafeImageFolder(root=str(splits['validate']), transform=tl_val_tf)
CLASS_NAMES = train_ds_tl.classes
print(f'Classes: {CLASS_NAMES} | Train: {len(train_ds_tl)} | Val: {len(val_ds_tl)}')

train_dl_tl = DataLoader(train_ds_tl, batch_size=32, shuffle=True, num_workers=4)
val_dl_tl = DataLoader(val_ds_tl, batch_size=32, shuffle=False, num_workers=4)

```

```
Classes: ['edible', 'poisonous'] | Train: 2,256 | Val: 282
```

```
In [13]: # — Load ResNet18 with correct modern API (pretrained=True is deprecated) ——  
resnet = models.resnet18(weights=models.ResNet18_Weights.DEFAULT)  
  
for param in resnet.parameters():          # freeze entire backbone  
    param.requires_grad = False  
  
# Replace head: 512 features -> dropout -> 256 -> 2 classes  
resnet.fc = nn.Sequential(  
    nn.Dropout(0.4),  
    nn.Linear(resnet.fc.in_features, 256), nn.ReLU(),  
    nn.Dropout(0.3),  
    nn.Linear(256, 2)  
)  
resnet = resnet.to(DEVICE)  
  
trainable = sum(p.numel() for p in resnet.parameters() if p.requires_grad)  
total      = sum(p.numel() for p in resnet.parameters())  
print(f'Initially trainable: {trainable:,} / {total:,} ({100*trainable/total:.1f}%)
```

```
Initially trainable: 131,842 / 11,308,354 (1.2%)
```


```
In [14]: # — Train: Phase A (head only) then auto Phase B (Layer4 + head) ——  
FINETUNE_EPOCH = 8    # at this epoch, Layer4 unfreezes with differential LRs  
  
criterion_tl = nn.CrossEntropyLoss()  
optimizer_tl = optim.Adam(filter(lambda p: p.requires_grad, resnet.parameters()))  
scheduler_tl = optim.lr_scheduler.StepLR(optimizer_tl, step_size=5, gamma=0.3)  
  
resnet, history_tl, best_preds, best_labels = train_model(  
    resnet, train_dl_tl, val_dl_tl,  
    criterion_tl, optimizer_tl, scheduler_tl,  
    num_epochs=15,  
    finetune_epoch=FINETUNE_EPOCH,      # auto-switch inside the loop  
    finetune_layers=('layer4', 'fc'),   # what to unfreeze  
    finetune_lr=1e-4                  # backbone LR; head gets 1e-3  
)  
  
plot_curves(history_tl,  
            title='ResNet18 Transfer Learning – Training Curves',  
            finetune_epoch=FINETUNE_EPOCH)  
  
best_val = max(history_tl['val_acc'])  
print(f'ResNet18 – Best Val Accuracy: {best_val:.2f}%')  
if best_val >= 85:  
    print('✅ SUCCESS: Validation accuracy exceeds 85% target!')  
else:  
    print(f'⚠️ {best_val:.2f}% is below 85%. Try: ')  
    print('• Increase num_epochs to 20')  
    print('• Lower finetune_epoch to 5 (unfreeze earlier)')  
    print('• Add transforms.RandomAffine(degrees=0, translate=(0.1, 0.1))')
```

```
Epoch 01/15 [Train]:  0%           | 0/71 [00:00<?, ?it/s]  
Epoch 01/15 [Val] :  0%           | 0/9 [00:00<?, ?it/s]  
Epoch 01/15 Train Loss: 0.6684  Acc: 60.02% | Val Loss: 0.5114  Acc: 75.89%  
    ★ New best model saved (Val Acc: 75.89%)  
Epoch 02/15 [Train]:  0%           | 0/71 [00:00<?, ?it/s]  
Epoch 02/15 [Val] :  0%           | 0/9 [00:00<?, ?it/s]  
Epoch 02/15 Train Loss: 0.5904  Acc: 68.93% | Val Loss: 0.4792  Acc: 76.24%  
    ★ New best model saved (Val Acc: 76.24%)
```

```
Epoch 03/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 03/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 03/15 Train Loss: 0.5791 Acc: 69.68% | Val Loss: 0.4808 Acc: 78.72%
★ New best model saved (Val Acc: 78.72%)
Epoch 04/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 04/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 04/15 Train Loss: 0.5607 Acc: 71.05% | Val Loss: 0.4369 Acc: 80.50%
★ New best model saved (Val Acc: 80.50%)
Epoch 05/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 05/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 05/15 Train Loss: 0.5644 Acc: 70.08% | Val Loss: 0.4520 Acc: 82.98%
★ New best model saved (Val Acc: 82.98%)
Epoch 06/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 06/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 06/15 Train Loss: 0.5489 Acc: 70.48% | Val Loss: 0.4364 Acc: 82.62%
Epoch 07/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 07/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 07/15 Train Loss: 0.5346 Acc: 71.68% | Val Loss: 0.4283 Acc: 80.50%
Epoch 08/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 08/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 08/15 Train Loss: 0.5390 Acc: 71.41% | Val Loss: 0.4230 Acc: 82.27%
```

```
=====
Epoch 9: SWITCHING TO FINE-TUNE – unfreezing ('layer4', 'fc')
=====

Epoch 09/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 09/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 09/15 Train Loss: 0.4993 Acc: 75.09% | Val Loss: 0.3244 Acc: 85.82%
★ New best model saved (Val Acc: 85.82%)
Epoch 10/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 10/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 10/15 Train Loss: 0.3529 Acc: 83.47% | Val Loss: 0.3356 Acc: 84.04%
Epoch 11/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 11/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 11/15 Train Loss: 0.2730 Acc: 88.39% | Val Loss: 0.2245 Acc: 89.72%
★ New best model saved (Val Acc: 89.72%)
Epoch 12/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 12/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 12/15 Train Loss: 0.2186 Acc: 90.20% | Val Loss: 0.1633 Acc: 92.91%
★ New best model saved (Val Acc: 92.91%)
Epoch 13/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 13/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 13/15 Train Loss: 0.1691 Acc: 93.40% | Val Loss: 0.1951 Acc: 89.72%
Epoch 14/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 14/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 14/15 Train Loss: 0.1576 Acc: 94.50% | Val Loss: 0.1566 Acc: 93.26%
★ New best model saved (Val Acc: 93.26%)
Epoch 15/15 [Train]: 0% | 0/71 [00:00<?, ?it/s]
Epoch 15/15 [Val] : 0% | 0/9 [00:00<?, ?it/s]
Epoch 15/15 Train Loss: 0.1333 Acc: 94.73% | Val Loss: 0.1155 Acc: 95.39%
★ New best model saved (Val Acc: 95.39%)
```




ResNet18 – Best Val Accuracy: 95.39%

SUCCESS: Validation accuracy exceeds 85% target!

5. Evaluation — Confusion Matrix & Classification Report

```
In [15]: cm = confusion_matrix(best_labels, best_preds)
plt.figure(figsize=(7, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=CLASS_NAMES, yticklabels=CLASS_NAMES)
plt.title('Confusion Matrix — ResNet18 (Best Epoch)')
plt.ylabel('True Label')
plt.xlabel('Predicted Label')
plt.tight_layout()
plt.show()

print(classification_report(best_labels, best_preds, target_names=CLASS_NAMES))
```



	precision	recall	f1-score	support
edible	0.97	0.95	0.96	150
poisonous	0.94	0.96	0.95	132
accuracy			0.95	282
macro avg	0.95	0.95	0.95	282
weighted avg	0.95	0.95	0.95	282

6. Discussion & Model Comparison

```
In [16]: print('*'*55)
print('MODEL COMPARISON')
print('*'*55)
print(f'Baseline CNN Best Val Acc : {max(history_base["val_acc"]):.2f}%')
print(f'ResNet18 TL Best Val Acc : {max(history_tl["val_acc"]):.2f}%')
print()
print('Key takeaways:')
print('• Baseline CNN proves the problem is learnable (60-70%)')
print('• ResNet18 already knows edges/textures from 1.2M ImageNet images')
print('• Phase A (frozen backbone) = fast early gains with low risk')
print('• Phase B (unfreeze layer4) = fine-tunes mushroom-specific features')
print('• Differential LRs protect pretrained weights during fine-tuning')
print('• Heavy augmentation (flip, rotate, color jitter, crop) reduces overfitting')
print('• copy.deepcopy checkpointing ensures we return the best epoch, not the last')

# Save model
torch.save(resnet.state_dict(), 'resnet18_mushroom_best.pth')
print('\nModel saved -> resnet18_mushroom_best.pth')
```

```
=====
MODEL COMPARISON
=====
Baseline CNN Best Val Acc : 72.70%
ResNet18 TL Best Val Acc : 95.39%
```

Key takeaways:

- Baseline CNN proves the problem is learnable (60-70%)
- ResNet18 already knows edges/textures from 1.2M ImageNet images
- Phase A (frozen backbone) = fast early gains with low risk
- Phase B (unfreeze layer4) = fine-tunes mushroom-specific features
- Differential LRs protect pretrained weights during fine-tuning
- Heavy augmentation (flip, rotate, color jitter, crop) reduces overfitting
- copy.deepcopy checkpointing ensures we return the best epoch, not the last

Model saved -> resnet18_mushroom_best.pth

7. Final Evaluation on TEST SET

```
In [17]: # Use the SAME transforms as validation
test_ds = SafeImageFolder(
    root=str(splits['test']),
    transform=tl_val_tf
)

test_dl = DataLoader(test_ds, batch_size=32, shuffle=False, num_workers=0)
```

```

print(f"Test samples: {len(test_ds)}")
print(f"Classes: {test_ds.classes}")

# — Run inference ——————
model = resnet # already loaded with best weights from training
model.eval()

all_preds = []
all_labels = []

with torch.no_grad():
    for imgs, labels in tqdm(test_dl, desc="Testing"):
        imgs = imgs.to(DEVICE)
        labels = labels.to(DEVICE)

        outputs = model(imgs)
        _, preds = outputs.max(1)

        all_preds.extend(preds.cpu().numpy())
        all_labels.extend(labels.cpu().numpy())

# — Metrics ——————
cm = confusion_matrix(all_labels, all_preds)
print("\nClassification Report:")
print(classification_report(all_labels, all_preds, target_names=CLASS_NAMES))

# — Confusion Matrix Plot ——————
plt.figure(figsize=(7, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Greens',
            xticklabels=CLASS_NAMES, yticklabels=CLASS_NAMES)
plt.title("Confusion Matrix – TEST SET")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.tight_layout()
plt.show()

```

Test samples: 282
 Classes: ['edible', 'poisonous']
 Testing: 0%| 0/9 [00:00<?, ?it/s]
 Classification Report:

	precision	recall	f1-score	support
edible	0.93	0.95	0.94	150
poisonous	0.95	0.92	0.93	132
accuracy			0.94	282
macro avg	0.94	0.94	0.94	282
weighted avg	0.94	0.94	0.94	282

Confusion Matrix — TEST SET

